

RocksDB Storage Engine

Igor Canadi | Facebook



Overview

- Story of RocksDB
- Architecture
- Performance tuning
- Next steps



Story of RocksDB



PERCONA
LIVE

Pre-2011

- FB infrastructure – many custom-built key-value stores
- LevelDB released

Experimentation (2011 – 2013)

- First use-cases
- Not designed for server – many bottlenecks, stalls
- Optimization
- New features



Explosion (2013 – 2015)

- Open sourced RocksDB
- Big success within Facebook
- External traction – Linkedin, Yahoo, CockroachDB, ...



New Challenges (2015 -)

- Bring RocksDB to databases



mongoDB



PERCONA
LIVE

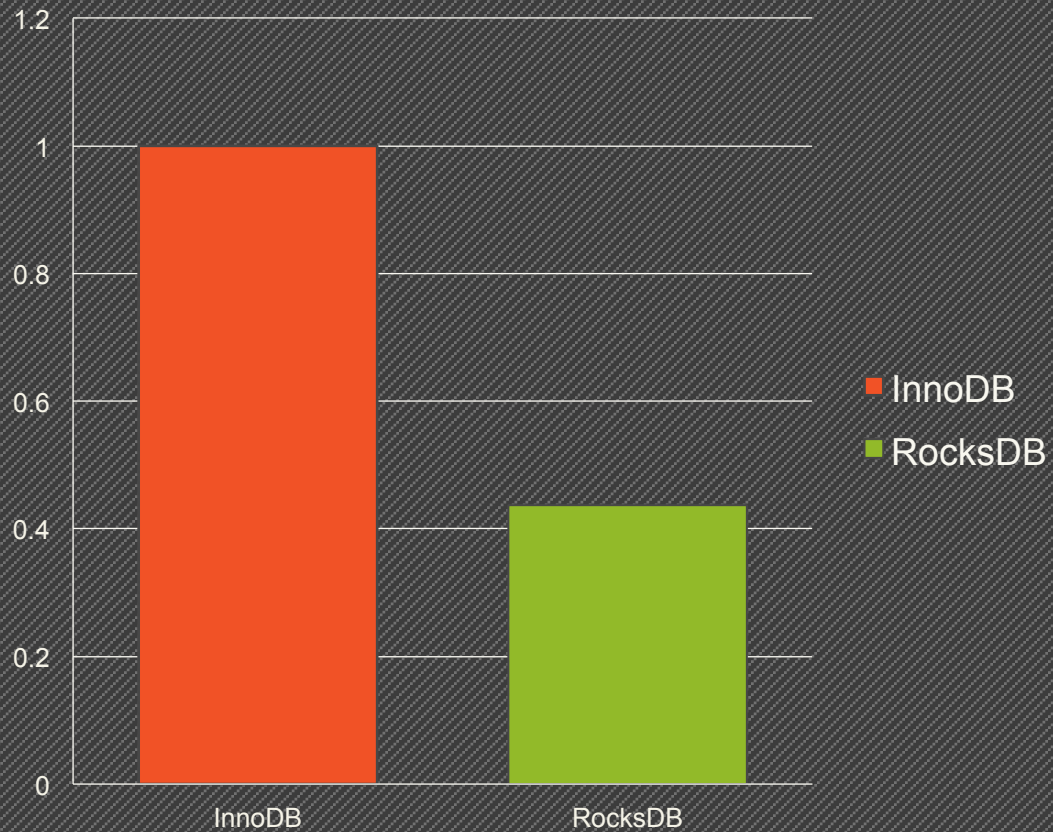
MongoRocks

- Running in production at Parse for 6 months
- Huge storage savings (5TB → 285GB)
- Document-level locking

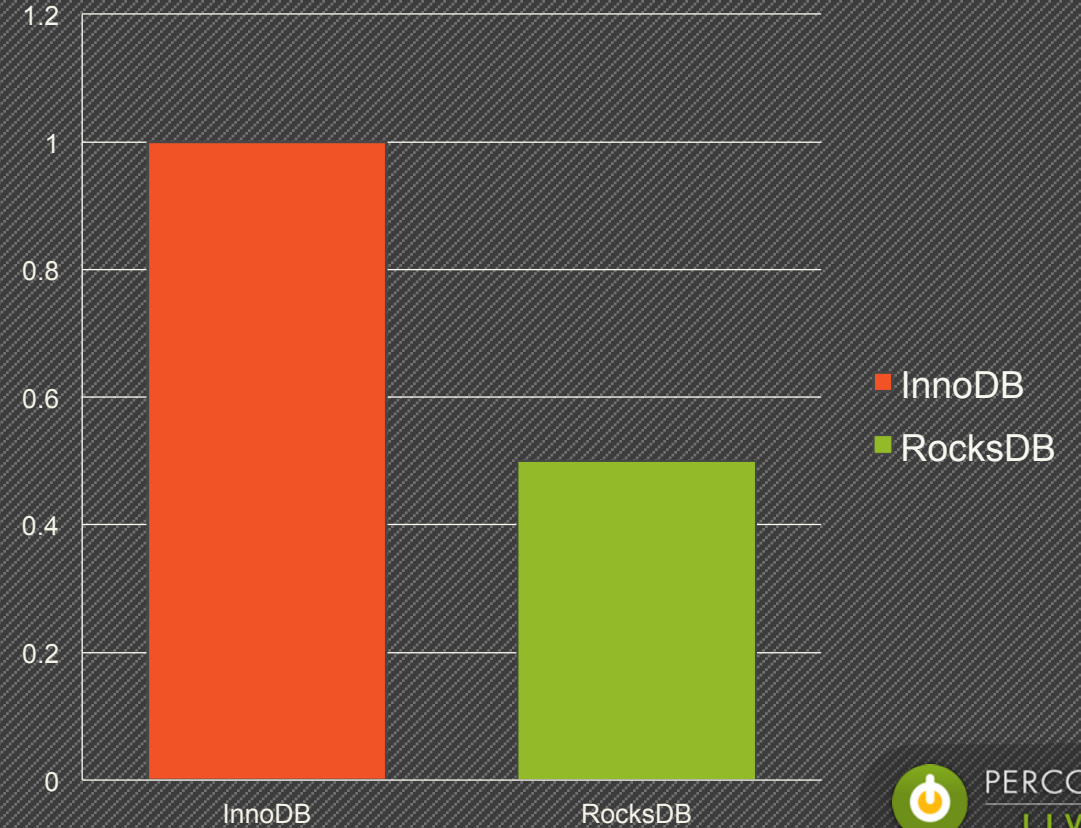


MyRocks

Database size (relative)



Bytes written (relative)




Architecture

Log Structured Merge Trees




PERCONA
LIVE


Log Structured Merge Trees




Log Structured Merge Trees – write



Log Structured Merge Trees – flush



Log Structured Merge Trees – compaction




Writes

- Foreground:
 - Writes go to memtable (skiplist) + write-ahead log
- Background:
 - When memtable is full, we flush to Level 0
 - When a level is full, we run compaction



Reads



Reads

- Point queries
 - Bloom filters reduce reads from storage
 - Usually only 1 read IO
- Range scans
 - Bloom filters don't help
 - Depends on amount of memory, 1-2 IO



RocksDB Files

```
rocksdb/> ls
MANIFEST-000032
000024.log
000031.log
000025.sst
000028.sst
000029.sst
000033.sst
000034.sst
LOG
LOG.old.1441234029851978
...
```



RocksDB Files – MANIFEST

- Atomical updates to database metadata

(initial state) Add file 1 Add file 2 Add file 3 Add file 4 ...	(flush) Add file 9 Mark log 6 persisted	(compaction) Add file 10 Add file 11 Remove file 9 Remove file 8	Add new column family “system”
--	---	--	-----------------------------------



RocksDB Files – Write-ahead log

- Persisted memtable state

Write (A, B)	Write (C, D) Write (E, F)	Delete(A)	Write(X, Y) Delete(C)
--------------	------------------------------	-----------	--------------------------



RocksDB Files – Table files

(Data block) <ul style="list-style-type: none">• compressed• prefix encoded	(Data block) <key, value>	(Data block)	(Data block)
(Data block)	(Data block)	(Data block)	(Data block)
(Index block) <key, block>	(Filter block)	(Statistics)	(Meta index block) Pointers to blocks



RocksDB Files – LOG files

- Debugging output
- Tuning options
- Information about flushes and compactions
- Performance statistics



Backups

- Table files are immutable
- Other files are append-only
- Easy and fast incremental backups
- Open sourced Rocks-Strata



Performance tuning



PERCONA
LIVE

Tombstones

- Deletions are deferred
- May cause higher P99 latencies
- Be careful with pathological workloads, e.g. queues



Caching

Block cache

- Managed by RocksDB
- Uncompressed data
- Defaults to 1/3 of RAM

Page cache

- Managed by kernel
- Compressed data



Memory usage

- Block cache
- Index and filter blocks (0.5 – 2% of the database)
- Memtables
- Blocks pinned by iterators



Reduce memory usage

- Reduce block cache size – will increase CPU
- Increase block size – decrease index size
- Turn off bloom filters on bottom level



Reduce CPU

- Profile the CPU usage
- Increase block cache size – will increase memory usage
- Turn off compression
- It might be tombstones



Reduce write amplification

- Write amplification = $5 * \text{num_levels}$
- Increase memtable and level 1 size
- Stronger (zlib, zstd) compression for bottom levels
- Try universal compaction



Next steps



PERCONA
LIVE

Next steps

- Increase performance & stability
- Deploy MyRocks at Facebook
- External adoption of MyRocks and MongoRocks
- Build an ecosystem



Thank you

